

Activity 14 - Stats Reasoning

Jocelyn Rojas & Seleona Hunter

Loading Packages

```
library(tidyverse) # For data wrangling
```

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr     1.1.4     v readr     2.1.6
v forcats   1.0.1     v stringr   1.6.0
v ggplot2   4.0.1     v tibble    3.3.1
v lubridate 1.9.4     v tidyr    1.3.2
v purrr    1.2.1

-- Conflicts -----
x dplyr::filter() masks stats::filter()
x dplyr::lag()    masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become non-conflicting.
```

```
library(brms) # For stats
```

```
Loading required package: Rcpp
Loading 'brms' package (version 2.23.0). Useful instructions
can be found by typing help('brms'). A more detailed introduction
to the package is available through vignette('brms_overview').
```

```
Attaching package: 'brms'
```

```
The following object is masked from 'package:stats':
```

```
ar
```

```
library(ggeffects) # for plotting model predictions
# Note: I needed to also install the `insight` and `see` packages to get `modelbased` to install
# install.packages('modelbased') # if you need to install this package
library(modelbased) # for plotting model predictions. supports the link scale (ggeffects does
```

```
Attaching package: 'modelbased'
```

```
The following objects are masked from 'package:ggeffects':
```

```
collapse_by_group, pool_predictions, residualize_over_grid
```

```
# install.packages('faraway') # if you need to install this package
library(faraway) # For data on galapagos species richness
```

```
Attaching package: 'faraway'
```

```
The following object is masked from 'package:brms':
```

```
epilepsy
```

Q1.1a and Q1.1b

1. Counts of Clarkia flowers in a meadow
 - Values of response variables: Positive real numbers (can't have a negative number of flowers), could include zeros (there may not be any flowers)
 - Distribution: Normal distribution
2. Whether or not a female elephant seal gives birth
 - Values of response variables: 0s and 1s (for yes: birth or no: did not give birth)
 - Distribution: Poisson distribution
3. The percent cover of red algae in the intertidal
 - Values of response variables: fractions, positive numbers
 - Distribution: Binomial distribution
4. Growth of a tree from one year to the next

- Values of response variables: positive integers
- Distribution: Normal distribution

5. The spatial area of a forest in square meters

- Values of response variables: integers, fractions
- Distribution: Binomial distribution

Q1.2

1. Response variable for final project: rate of flock movement (through a forest)
2. Values of response variables: positive real numbers, integers
3. Distribution: Normal distribution

1.2 GLM with a log link

```
# Read in the pre-stored data
data("gala")
# Check out the first 6 rows
head(gala)
```

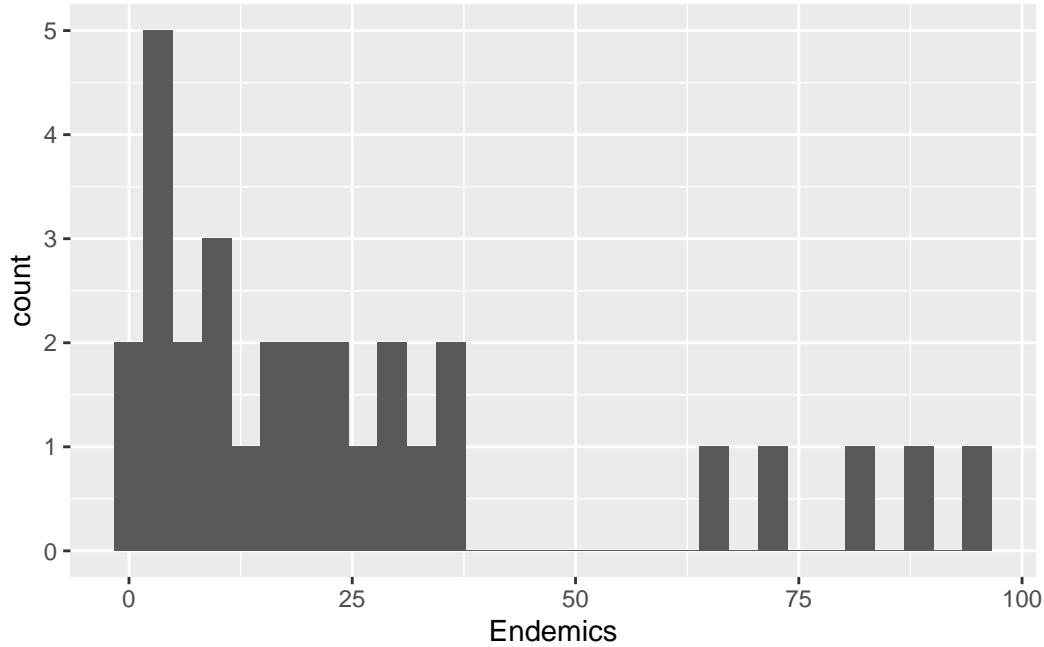
	Species	Endemics	Area	Elevation	Nearest	Scruz	Adjacent
Baltra	58	23	25.09	346	0.6	0.6	1.84
Bartolome	31	21	1.24	109	0.6	26.3	572.33
Caldwell	3	3	0.21	114	2.8	58.7	0.78
Champion	25	9	0.10	46	1.9	47.4	0.18
Coamano	2	1	0.05	77	1.9	1.9	903.82
Daphne.Major	18	11	0.34	119	8.0	8.0	1.84

```
# Look at the help page too!
?gala
```

Q1.3 Plot a histogram of the response variable Endemics

```
ggplot(data = gala, aes(x=Endemics)) +
  geom_histogram()

`stat_bin()` using `bins = 30`. Pick better value `binwidth`.
```

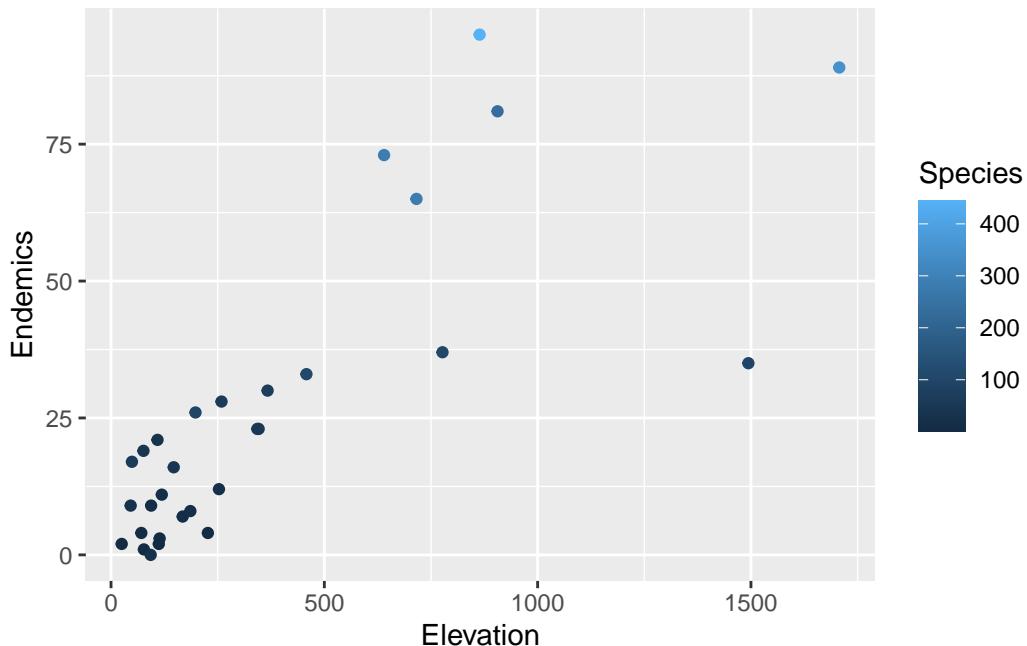


Q1.3 Answer

This data does not appear to be normally distributed. The data seems to be skewed to the left with outliers far to the right in the graph.

Q1.4 Plot Endemics ~ Elevation

```
ggplot(data = gala, aes(x = Elevation, y = Endemics, color = Species)) +  
  geom_point()
```



```
labs(
  x = "Elevation",
  y = "Endemics"
)
```

```
<ggplot2::labels> List of 2
$ x: chr "Elevation"
$ y: chr "Endemics"
```

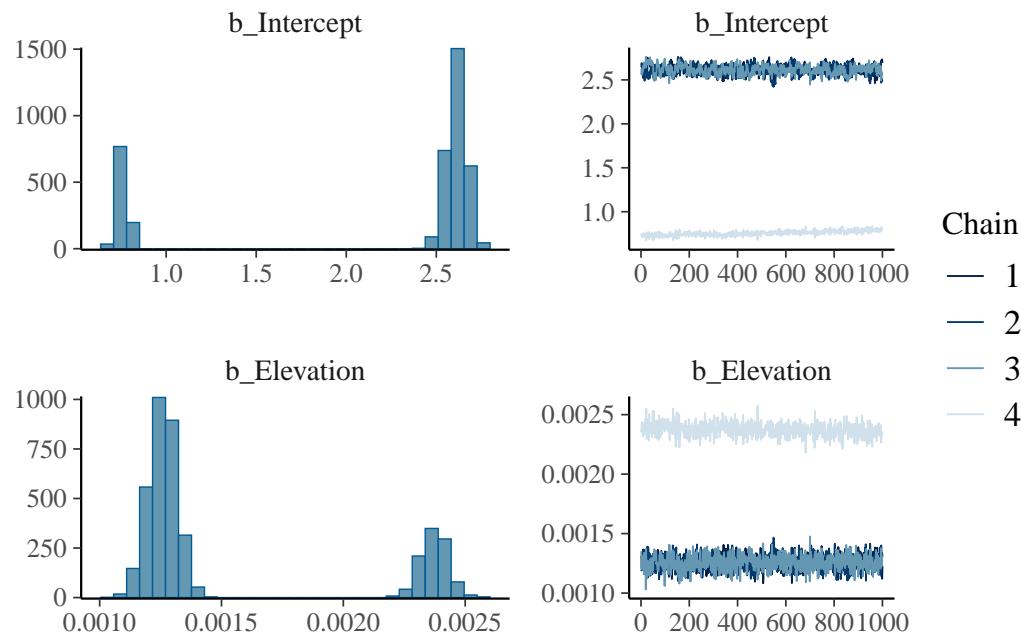
Run the model

```
# Endemics ~ Elevation
m.elev <-
  brm(data = gala, # Give the model the penguins data
       # Choose a poisson distribution - THIS IS THE NEW PART!
       family = poisson(link = "log"),
       # Specify the model here.
       Endemics ~ 1 + Elevation,
       # Here's where you specify parameters for executing the Markov chains
       # We're using similar to the defaults, except we set cores to 4 so the analysis runs faster
       iter = 2000, warmup = 1000, chains = 4, cores = 4,
```

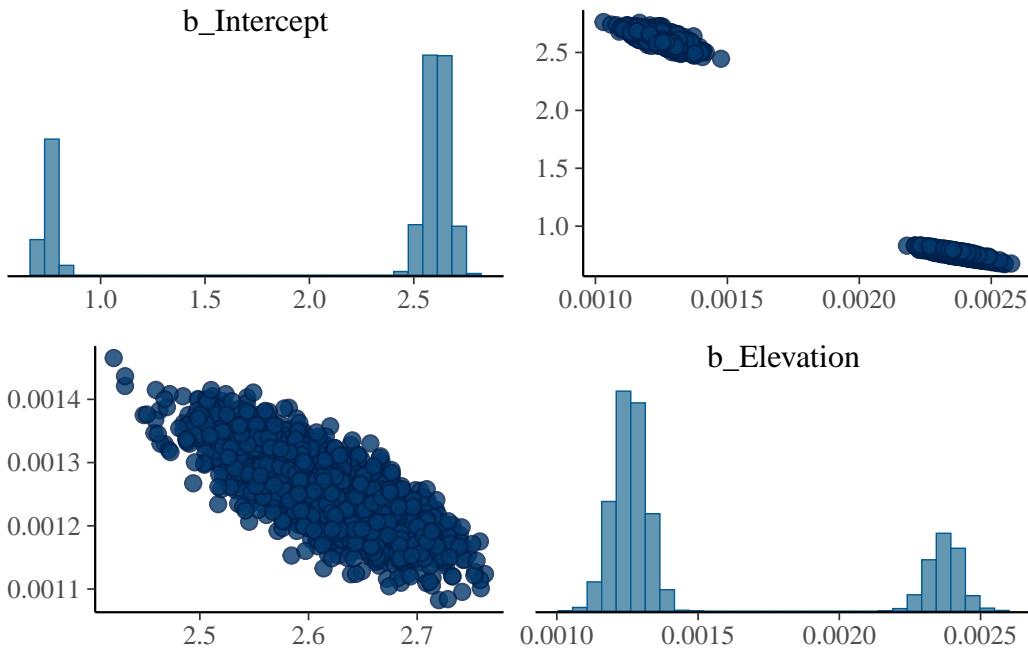
```
# Save the fitted model object as output - helpful for reloading in the output later  
file = "output/m.elev")
```

model summary

```
plot(m.elev)
```



```
pairs(m.elev)
```



```
summary(m.elev)
```

Warning: Parts of the model have not converged (some Rhats are > 1.05). Be careful when analysing the results! We recommend running more iterations and/or setting stronger priors.

```

Family: poisson
Links: mu = log
Formula: Endemics ~ 1 + Elevation
Data: gala (Number of observations: 30)
Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
      total post-warmup draws = 4000

```

Regression Coefficients:

	Estimate	Est.Error	1-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	2.15	0.81	0.72	2.71	1.57	7	13
Elevation	0.00	0.00	0.00	0.00	1.53	7	27

Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS and Tail_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

Q1.5 Evaluate the output

Based on the summary output, the model fitting algorithm did not converge. Our RHat value exceeds 1 and the chains do not align.

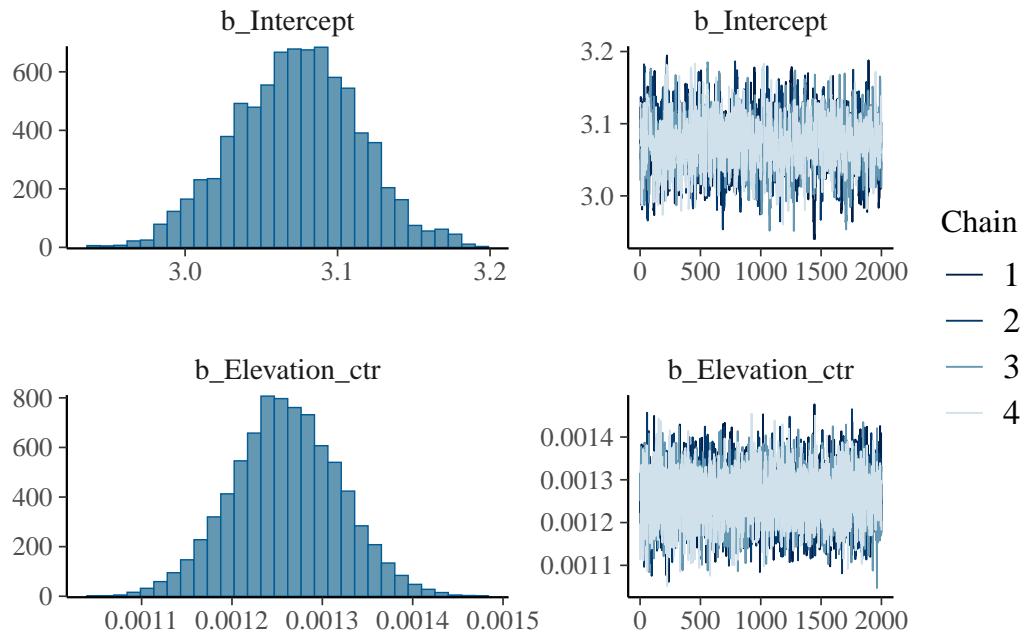
Q1.6 Center the predictors

```
gala <- gala |>
  mutate(Elevation_ctr = Elevation - mean(Elevation))

# Endemics ~ Elevation
m.elev2 <-
  brm(data = gala, # Give the model the penguins data
       # Choose a poisson distribution - THIS IS THE NEW PART!
       family = poisson(link = "log"),
       # Specify the model here.
       Endemics ~ 1 + Elevation_ctr,
       # Here's where you specify parameters for executing the Markov chains
       # We're using similar to the defaults, except we set cores to 4 so the analysis runs faster
       iter = 6000, warmup = 4000, chains = 4, cores = 4,
       prior = prior(normal(0, 0.1), class = b),
       # Save the fitted model object as output - helpful for reloading in the output later
       file = "output/m.elev2")
```

checking the model

```
plot(m.elev2)
```



```
summary(m.elev2)
```

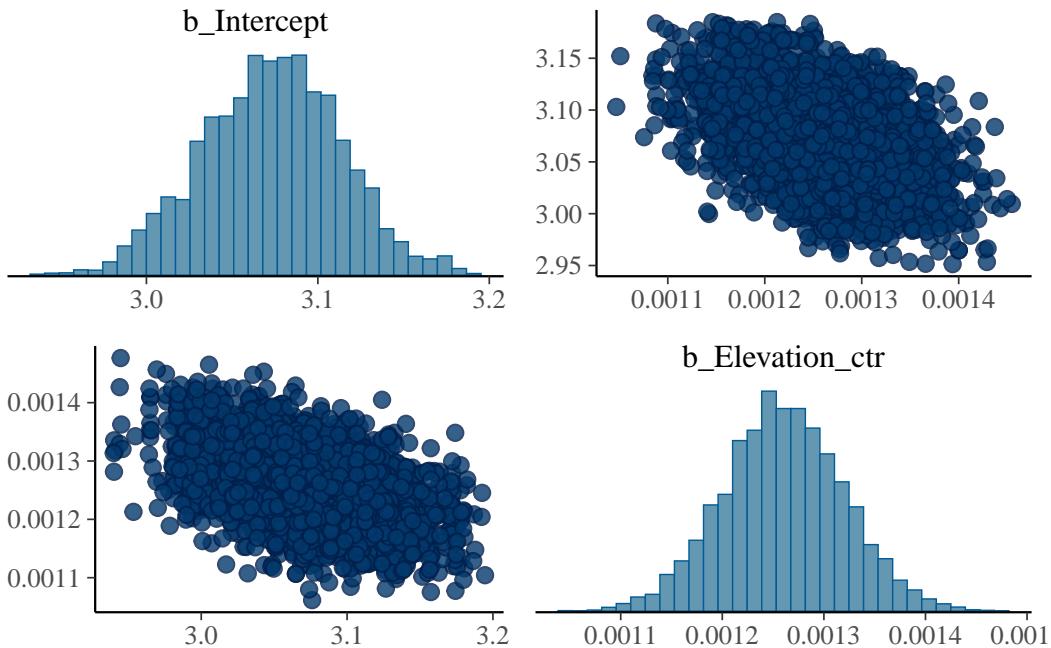
Family: poisson
 Links: mu = log
 Formula: Endemics ~ 1 + Elevation_ctr
 Data: gala (Number of observations: 30)
 Draws: 4 chains, each with iter = 6000; warmup = 4000; thin = 1;
 total post-warmup draws = 8000

Regression Coefficients:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	3.07	0.04	2.99	3.15	1.00	1137	1217
Elevation_ctr	0.00	0.00	0.00	0.00	1.00	3801	4656

Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS and Tail_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

```
pairs(m.elev2)
```



```
# summary(m.elev2)
print(m.elev2, digits = 4) #alternative summary with more decimals
```

```
Family: poisson
Links: mu = log
Formula: Endemics ~ 1 + Elevation_ctr
Data: gala (Number of observations: 30)
Draws: 4 chains, each with iter = 6000; warmup = 4000; thin = 1;
      total post-warmup draws = 8000
```

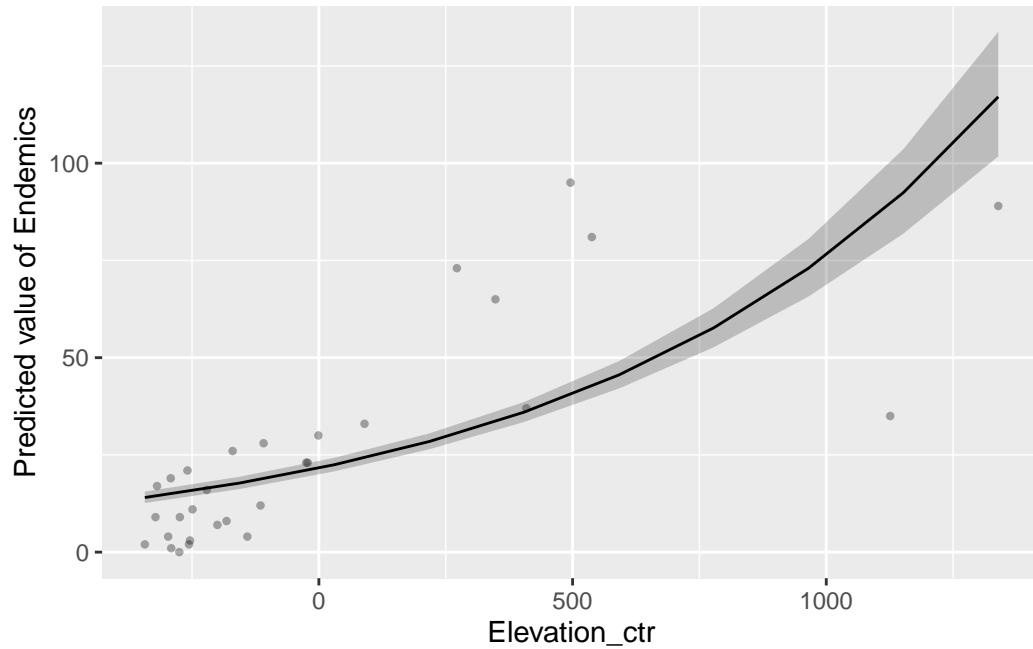
Regression Coefficients:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	3.0731	0.0404	2.9936	3.1513	1.0023	1137	1217
Elevation_ctr	0.0013	0.0001	0.0011	0.0014	1.0001	3801	4656

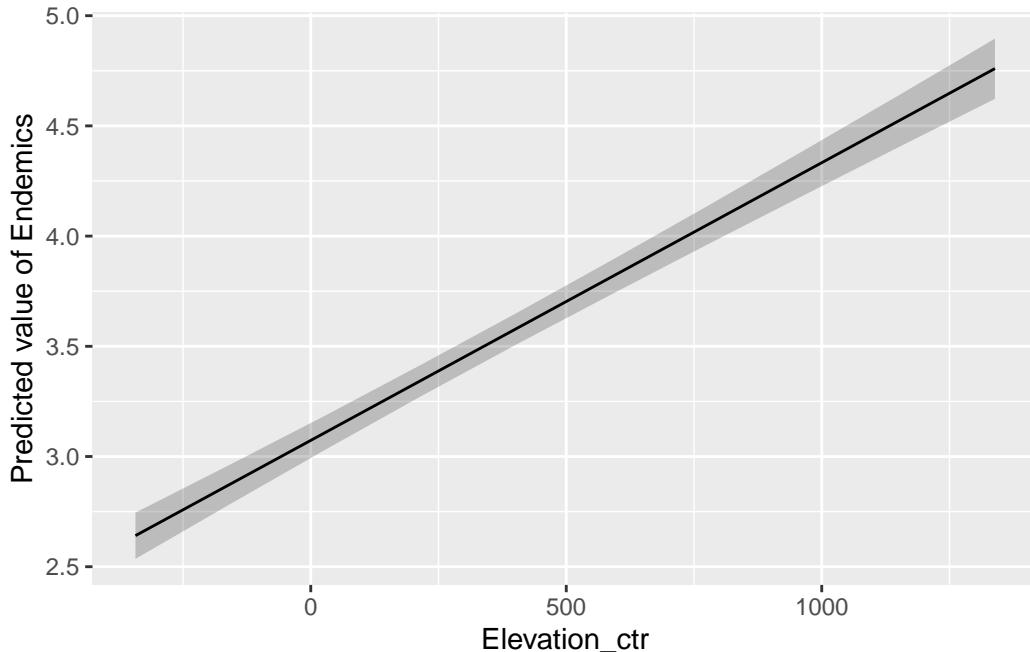
Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS and Tail_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

Plot the posterior

```
#plotting posterior predictions
preds <- estimate_expectation(m.elev2, by = 'Elevation_ctr')
plot(preds, show_data = TRUE)
```



```
# visualizing linear model
predslog <- estimate_expectation(m.elev2, by = 'Elevation_ctr', predict = 'link')
plot(predslog)
```



```
##Interpreting Link scale coeffs
```

```
print(m.elev2, digits = 4)
```

```
Family: poisson
Links: mu = log
Formula: Endemics ~ 1 + Elevation_ctr
Data: gala (Number of observations: 30)
Draws: 4 chains, each with iter = 6000; warmup = 4000; thin = 1;
      total post-warmup draws = 8000
```

Regression Coefficients:

	Estimate	Est.Error	1-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	3.0731	0.0404	2.9936	3.1513	1.0023	1137	1217
Elevation_ctr	0.0013	0.0001	0.0011	0.0014	1.0001	3801	4656

Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS and Tail_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

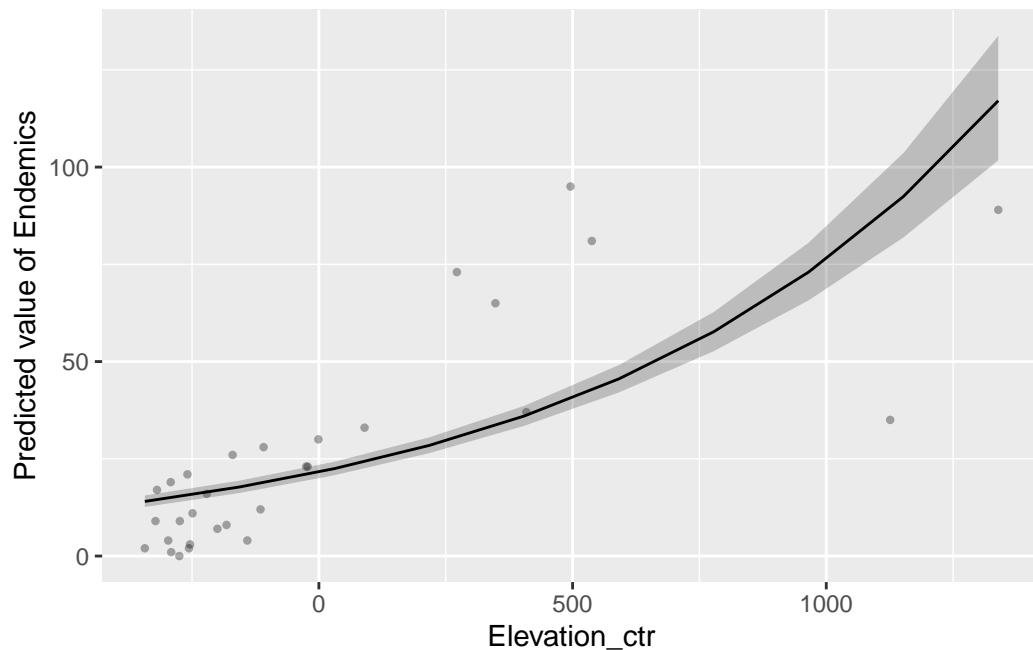
1. Backtransform

```
# backtransforming slope value  
exp(0.0013)
```

```
[1] 1.001301
```

2. Interpret

```
plot(preds, show_data = TRUE)
```



```
# for every 1m of elevation, you get an increase in 1.0013 times as many species as the previous
```

##Q1.7 What is the percent change on the response scale?

1. Number of Clarkias blooming as a function of temperature in Celsius: 1.09

```
exp(1.09)
```

```
[1] 2.974274
```

```
(2.974274 - 1)*100
```

```
[1] 197.4274
```

```
# Back transformed: 197.4274% increase in number of Clarkias for every 1 degree increase in
```

2. Density of sea urchins per square meter in a quadrat as a function of number of sea otters: -2.5

```
(exp(-2.5) - 1)*100
```

```
[1] -91.7915
```

```
# Back transformed: 0.082085
```

```
# For every increase in 1 sea otter there is a 91.7915% decrease in density of sea urchins p
```

3. Number of tomatoes per plant as a function of kg of fertilizer: 6.24

```
exp(6.24)
```

```
[1] 512.8585
```

```
(exp(6.24)-1)*100
```

```
[1] 51185.85
```

```
# Back transformed: 512.8585 tomatoes per plant
```

```
# For every increase in 1kg of fertilizer, there will be a 51185.85% increase in tomatoes pe
```

DIY: Run a model of non-endemic species ~ distance from Santa Cruz Island

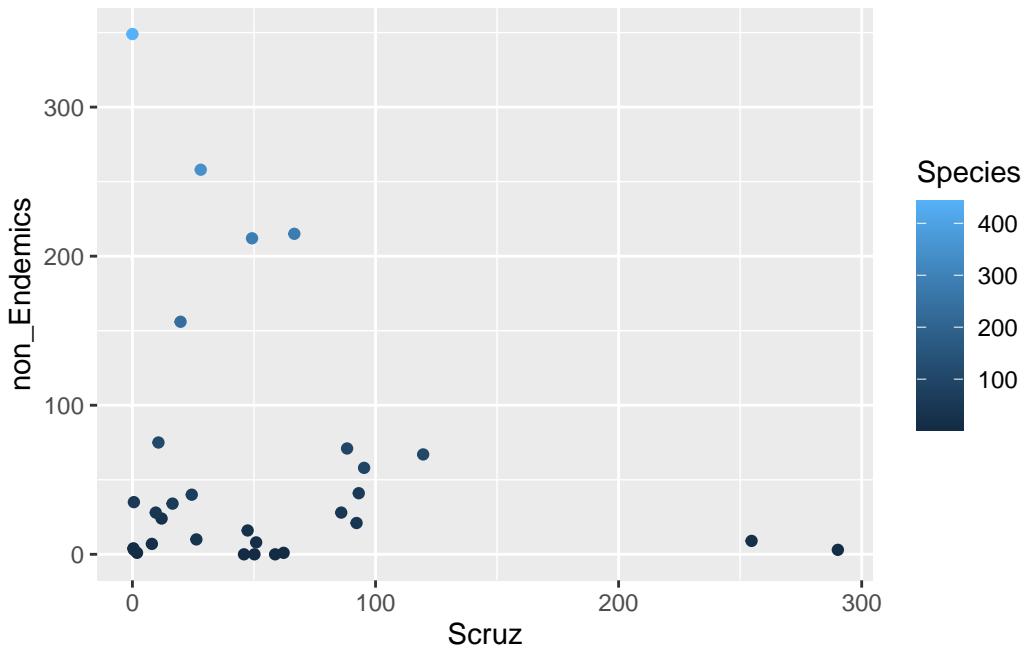
Q1.8 Create a non-endemic column

```

gala2 <- gala %>%
  mutate(non_Endemics = Species - Endemics)

# making a plot with new tibble
ggplot(data = gala2, aes(x = Scruz, y = non_Endemics, color = Species)) +
  geom_point()

```



```

  labs(
    x = "Santa Cruz Island (km)",
    y = "Non Endemic Species"
  )

```

```

<ggplot2::labels> List of 2
$ x: chr "Santa Cruz Island (km)"
$ y: chr "Non Endemic Species"

```

Q1.9 Run a model of non Endemics ~ distance from Santa Cruz Island

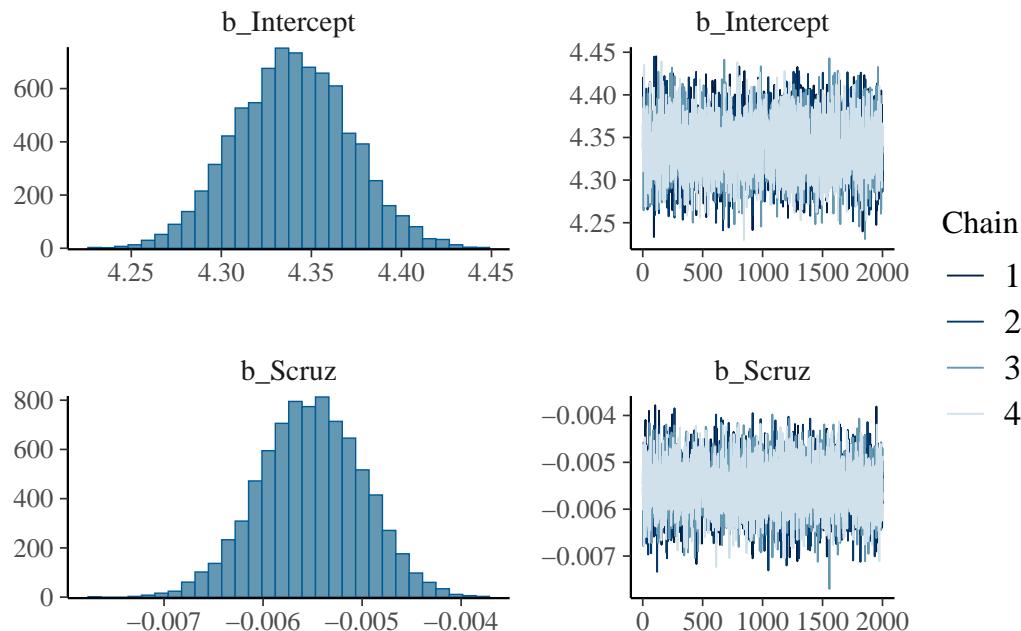
```

#model for non edemics as function of distance from santa cruz islands
gala2model <-
  brm(data = gala2, # Give the model the penguins data
       # Choose a poisson distribution - THIS IS THE NEW PART!
       family = poisson(link = "log"),
       # Specify the model here.
       non_Endemics ~ 1 + Scruz,
       # Here's where you specify parameters for executing the Markov chains
       # We're using similar to the defaults, except we set cores to 4 so the analysis runs faster
       iter = 6000, warmup = 4000, chains = 4, cores = 4,
       prior = prior(normal(0, 0.1), class = b),
       # Save the fitted model object as output - helpful for reloading in the output later
       file = "output/gala2model")

```

Q1.10 Evaluate the output

```
plot(gala2model)
```



```

print(gala2model, digits = 4)

Family: poisson
Links: mu = log
Formula: non_Endemics ~ 1 + Scruz
Data: gala2 (Number of observations: 30)
Draws: 4 chains, each with iter = 6000; warmup = 4000; thin = 1;
       total post-warmup draws = 8000

Regression Coefficients:
             Estimate Est.Error 1-95% CI u-95% CI   Rhat Bulk_ESS Tail_ESS
Intercept     4.3396    0.0316   4.2785   4.4020 1.0001      5774      5636
Scruz        -0.0055    0.0005  -0.0066  -0.0045 1.0006      5365      5369

```

Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS and Tail_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

Our model ran correctly. The Rhat is 1, the chains overlap, and the posteriors are normally distributed. Our confidence intervals do not contain zero.

Q1.11 Interpret the output

1. The number of non endemic species decreases as distance from Santa Cruz island increases.
 - a. Original output on log scale: 0.0055 log(non-endemic species)
 - b. Backtransformed: 0.9945151

```

exp(-0.0055)

```

[1] 0.9945151

- c. % change: For every 1 km increase in the distance from Santa Cruz island, there is a 0.54% decrease in non endemic species.

```

(exp(-0.0055)-1)*100

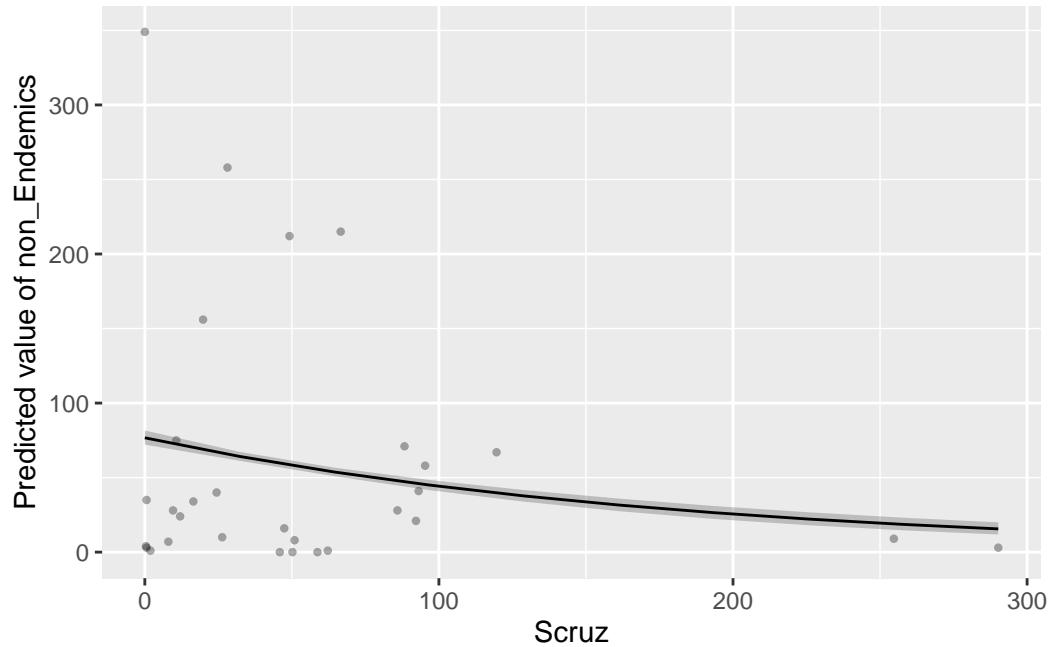
```

[1] -0.5484903

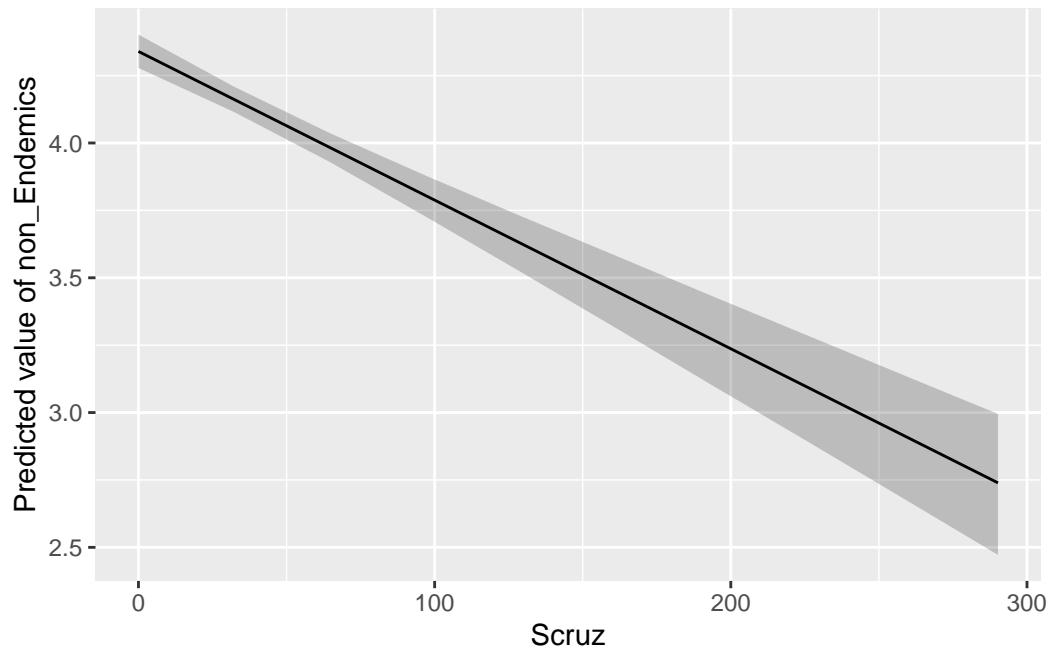
2. Yes it appears the slope estimate is different from zero because the confidence intervals do not contain zero.

Q1.12 Plot the posterior

```
# log link scale  
preds <- estimate_expectation(gala2model, by = 'Scruz')  
plot(preds, show_data = TRUE)
```



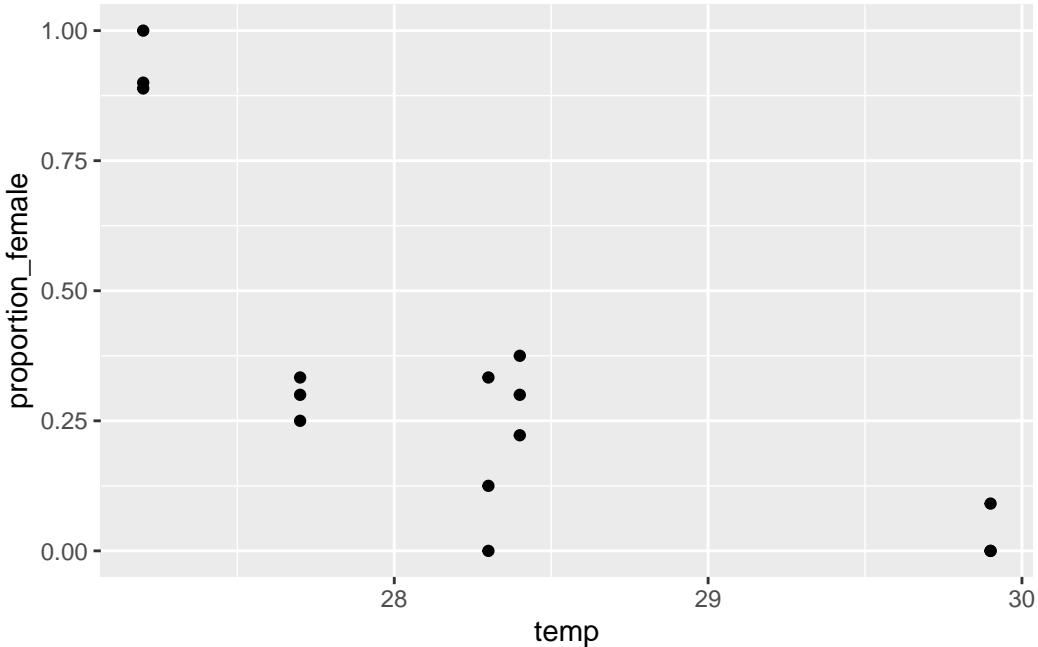
```
# visualizing linear model  
predslog <- estimate_expectation(gala2model, by = 'Scruz', predict = 'link')  
plot(predslog)
```



1.3 GLM with a logit link

```
# loading data
turtle <- faraway::turtle %>%
  mutate(total_turtles = male + female,
        proportion_female = female/total_turtles)

# plotting data
turtle %>%
  ggplot(aes(x = temp, y = proportion_female)) +
  geom_point()
```



```
# running model

m.turt <-
  brm(data = turtle, # Give the model the data
       # Choose a binomial distribution - THIS IS THE NEW PART!
       family = binomial(link = "logit"),
       # Specify the model here.
       female | trials(total_turtles) ~ 1 + temp,
       # Here's where you specify parameters for executing the Markov chains
       # We're using similar to the defaults, except we set cores to 4 so the analysis runs faster
       iter = 4000, warmup = 1000, chains = 4, cores = 4,
       # Save the fitted model object as output - helpful for reloading in the output later
       file = "output/m.turt")
```

```
# model output
summary(m.turt)
```

```
Family: binomial
Links: mu = logit
Formula: female | trials(total_turtles) ~ 1 + temp
Data: turtle (Number of observations: 15)
Draws: 4 chains, each with iter = 4000; warmup = 1000; thin = 1;
```

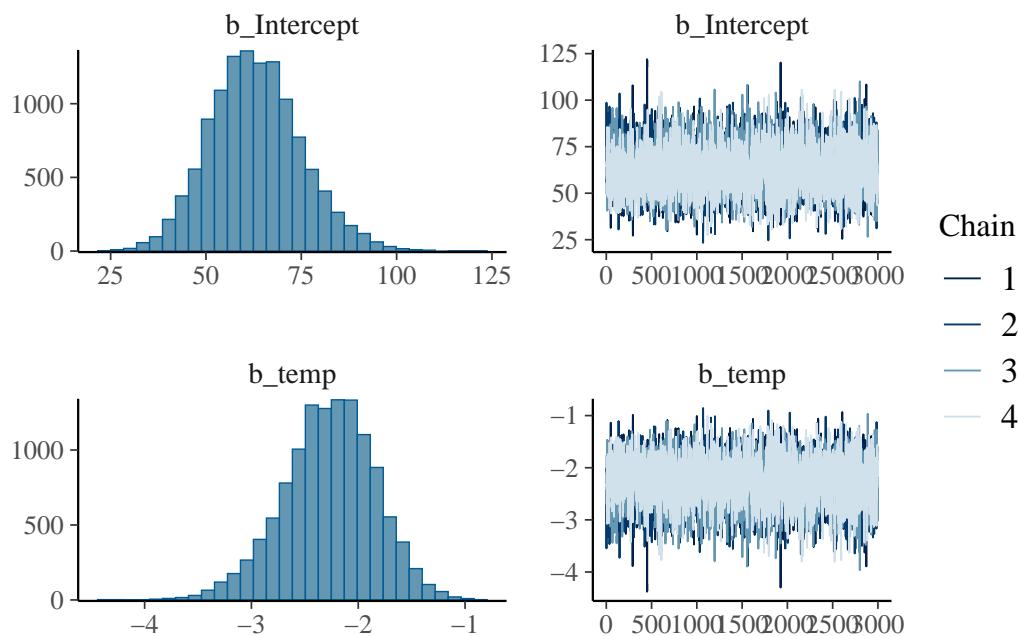
```
total post-warmup draws = 12000
```

Regression Coefficients:

	Estimate	Est.Error	1-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	63.00	12.08	40.68	88.62	1.00	4675	5550
temp	-2.27	0.43	-3.19	-1.47	1.00	4623	5557

Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS and Tail_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

```
plot(m.turt)
```



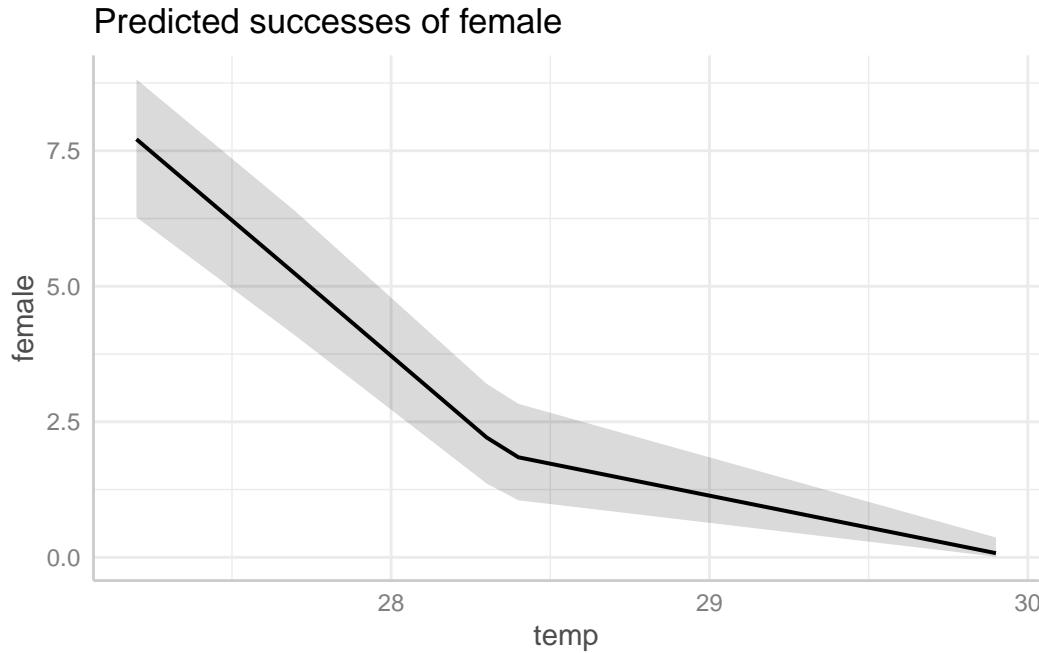
Model interpretations

Our model ran correctly. The Rhat is 1, posteriors are normally distributed, the chains overlap, and the confidence intervals do not include zero.

Interpreting plot predictions

```
pred <- predict_response(m.turt, condition = c(total_turtles = 10))

plot(pred)
```



2. Multilevel models

Q2.1 Fixed effects vs random effects

1. Student high school graduation rates as a function of: parental income, state of residence, and school district
 - Fixed: state of residence, school district
 - Random: parental income
2. Density of kelp as a function of: latitude, site, transect number, and density of sea urchins
 - Fixed: latitude, transect number
 - Random: density of sea urchins, site

3. Probability of whale giving birth as a function of: age, annual temperature, year, individual ID

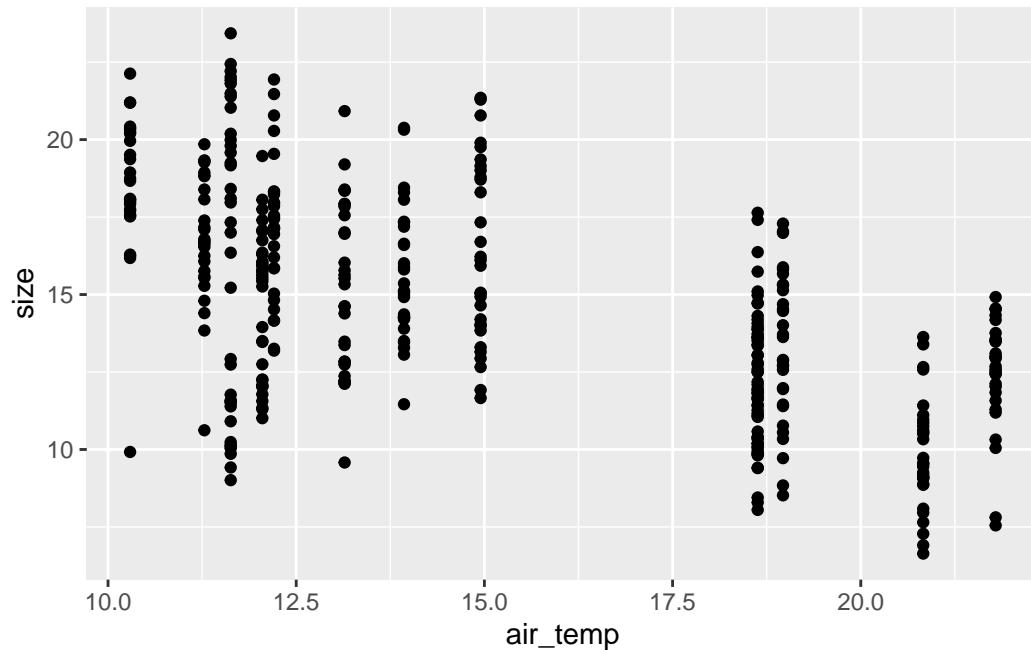
- Fixed: individual ID, year
- Random: age, annual temperature

Random Effects

```
# loading crab data

pie_crab <- lterdatasampler::pie_crab %>%
  mutate(site = as.factor(site))

pie_crab %>%
  ggplot(aes(x = air_temp, y = size)) +
  geom_point()
```



```
# running crab model
m.watertemp <-
  brm(data = pie_crab, # Give the model the penguins data
```

```

# Use a gamma distribution
family = Gamma(link = "log"),
# Specify the model here.
size ~ 1 + water_temp,
# Here's where you specify parameters for executing the Markov chains
# We're using similar to the defaults, except we set cores to 4 so the analysis runs faster
iter = 2000, warmup = 1000, chains = 4, cores = 4,
# Save the fitted model object as output - helpful for reloading in the output later
file = "output/m.watertemp")

print(m.watertemp, digits = 3)

```

Family: gamma
 Links: mu = log
 Formula: size ~ 1 + water_temp
 Data: pie_crab (Number of observations: 392)
 Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
 total post-warmup draws = 4000

Regression Coefficients:

	Estimate	Est.Error	1-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	3.356	0.055	3.250	3.465	1.001	5308	3289
water_temp	-0.039	0.003	-0.044	-0.033	1.001	5255	3218

Further Distributional Parameters:

	Estimate	Est.Error	1-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
shape	23.242	1.644	20.153	26.554	1.001	2572	2582

Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS and Tail_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

```

#site as a random effect

m.watertemp.site <-
  brm(data = pie_crab, # Give the model the penguins data
       # Use a gamma distribution
       family = Gamma(link = "log"),
       # Specify the model here.
       size ~ 1 + water_temp + (1|site),
       # Here's where you specify parameters for executing the Markov chains

```

```

# We're using similar to the defaults, except we set cores to 4 so the analysis runs faster
iter = 2000, warmup = 1000, chains = 4, cores = 4,
# Save the fitted model object as output - helpful for reloading in the output later
file = "output/m.watertemp.site")

print(m.watertemp.site, digits = 3)

```

Family: gamma

Links: mu = log

Formula: size ~ 1 + water_temp + (1 | site)

Data: pie_crab (Number of observations: 392)

Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
total post-warmup draws = 4000

Multilevel Hyperparameters:

~site (Number of levels: 13)

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
sd(Intercept)	0.125	0.035	0.077	0.213	1.002	886	1198

Regression Coefficients:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	3.372	0.203	2.970	3.778	1.003	1055	1133
water_temp	-0.040	0.011	-0.062	-0.017	1.003	1072	1220

Further Distributional Parameters:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
shape	30.115	2.168	26.002	34.555	1.000	3067	2243

Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS and Tail_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

Q2.2 What is the effect of water_temp on crab size on the response scale?

1. What is the effect of water_temp on crab size?
 - a) original output on the log scale: -0.039
 - b) your backtransformed value

```
exp(-0.039)
```

[1] 0.9617507

c) the percent change that this translates to. Describe the effect using the proper units.

For every 1 degree Celsius increase in water temperature, there crab size decrease by 3.82%.

```
(exp(-0.039)-1)*100
```

[1] -3.824929

2. Does it seem like the slope estimate is different from zero? Why?

Yes, the confidence interval for slope does not include zero.

Q2.3 Compare WAIC and PSIS of the two models

The model including water temperature and site was better predictive power than the model that only included water temperature.

```
loo(m.watertemp.site)
```

Computed from 4000 by 392 log-likelihood matrix.

	Estimate	SE
elpd_loo	-938.3	13.7
p_loo	12.1	0.8
looic	1876.6	27.4

MCSE of elpd_loo is 0.1.

MCSE and ESS estimates assume MCMC draws (r_eff in [0.6, 1.8]).

All Pareto k estimates are good (k < 0.7).

See help('pareto-k-diagnostic') for details.

```
loo(m.watertemp)
```

```
Computed from 4000 by 392 log-likelihood matrix.
```

	Estimate	SE
elpd_loo	-984.7	12.1
p_loo	2.8	0.2
looic	1969.5	24.3

MCSE of elpd_loo is 0.0.
MCSE and ESS estimates assume MCMC draws (r_eff in [0.6, 1.1]).

All Pareto k estimates are good (k < 0.7).
See help('pareto-k-diagnostic') for details.

```
waic(m.watertemp.site)
```

```
Computed from 4000 by 392 log-likelihood matrix.
```

	Estimate	SE
elpd_waic	-938.3	13.7
p_waic	12.0	0.8
waic	1876.6	27.4

```
waic(m.watertemp)
```

```
Computed from 4000 by 392 log-likelihood matrix.
```

	Estimate	SE
elpd_waic	-984.7	12.1
p_waic	2.8	0.2
waic	1969.5	24.3

Predict response of random effects

```
preds <- predict_response(m.watertemp.site,  
                           interval = "prediction",
```

```
    terms = "site",
    type = "random")  
  
plot(preds)
```

